

svhn_prediction

September 22, 2020

1 Import modules

```
[63]: import tensorflow as tf
assert tf.__version__ == '2.3.0', "tf version not 2.3.0"
import tensorflow_datasets as tfds
from tensorflow.keras import Sequential
from tensorflow.keras.layers import (Dense, Flatten, Conv2D, MaxPool2D,
                                     GlobalMaxPool2D, Dropout,
                                     BatchNormalization,
                                     Activation)

# import callbacks
from tensorflow.keras.callbacks import (ModelCheckpoint, EarlyStopping,
                                       TerminateOnNaN,
                                       ReduceLROnPlateau, TensorBoard,
                                       LearningRateScheduler)

%load_ext tensorboard

import numpy as np
import matplotlib.pyplot as plt
```

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

2 Download dataset from tfds

```
[64]: (train_ds, test_ds), info = tfds.load('svhn_cropped', split=['train', 'test'],
                                           with_info=True, shuffle_files=True,
                                           as_supervised=True)

train_ds, test_ds, info
```

```
[64]: (<DatasetV1Adapter shapes: ((32, 32, 3), ()), types: (tf.uint8, tf.int64)>,
      <DatasetV1Adapter shapes: ((32, 32, 3), ()), types: (tf.uint8, tf.int64)>,
      tfds.core.DatasetInfo(
        name='svhn_cropped',
        version=3.0.0,
        description='The Street View House Numbers (SVHN) Dataset is an image digit
```

```

recognition dataset of over 600,000 digit images coming from real world data.
Images are cropped to 32x32.',
    homepage='http://ufldl.stanford.edu/housenumbers/',
    features=FeaturesDict({
        'image': Image(shape=(32, 32, 3), dtype=tf.uint8),
        'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=10),
    }),
    total_num_examples=630420,
    splits={
        'extra': 531131,
        'test': 26032,
        'train': 73257,
    },
    supervised_keys=('image', 'label'),
    citation="""@article{Netzer2011,
    author = {Netzer, Yuval and Wang, Tao and Coates, Adam and Bissacco,
Alessandro and Wu, Bo and Ng, Andrew Y},
    booktitle = {Advances in Neural Information Processing Systems ({NIPS})},
    title = {Reading Digits in Natural Images with Unsupervised Feature
Learning},
    year = {2011}
}""",
    redistribution_info=,
))

```

3 Create major constants

```

[65]: cp_path_every_epoch = "checkpoint_every_epoch/"
cp_path_best_epoch = "checkpoint_best_epoch/"
log_dir = "logs_fit"

img_shape = tuple(train_ds.output_shapes[0])
img_shape

```

```
[65]: (32, 32, 3)
```

4 Data preview

```

[66]: # help(train_ds)
plt.figure(figsize=(8, 3.5))

np_ds = tfds.as_numpy(train_ds)
for (image, label), ii in zip(np_ds, [i for i in range(10)]):
    plt.subplot(2, 5, ii + 1)
    plt.imshow(image)

```

```

plt.xticks([])
plt.yticks([])
plt.title(label)
plt.show()

plt.figure(figsize=(8, 3.5))
for (image, label), ii in zip(np_ds, [i for i in range(10)]):
    plt.subplot(2, 5, ii + 1)
    plt.imshow(image.mean(-1), cmap='gray')
    plt.xticks([])
    plt.yticks([])
    plt.title(label)
plt.show()

```



5 Data preprocessing

```
[67]: def normalize_img(image, label):  
        """Normalizes images: `uint8` -> `float32`. """  
        return tf.cast(image, tf.float32) / 255., label  
  
print(train_ds)  
ds_train = train_ds.map(  
    normalize_img, num_parallel_calls=tf.data.experimental.AUTOTUNE)  
ds_train = ds_train.cache()  
ds_train = ds_train.batch(128)  
ds_train = ds_train.prefetch(tf.data.experimental.AUTOTUNE)  
print(ds_train)  
  
<DatasetV1Adapter shapes: ((32, 32, 3), ()), types: (tf.uint8, tf.int64)>  
<DatasetV1Adapter shapes: ((None, 32, 32, 3), (None,)), types: (tf.float32,  
tf.int64)>
```

```
[68]: print(test_ds)  
ds_test = test_ds.map(  
    normalize_img, num_parallel_calls=tf.data.experimental.AUTOTUNE)  
ds_test = ds_test.cache()  
ds_test = ds_test.batch(128)  
ds_test = ds_test.prefetch(tf.data.experimental.AUTOTUNE)  
print(ds_test)  
  
<DatasetV1Adapter shapes: ((32, 32, 3), ()), types: (tf.uint8, tf.int64)>  
<DatasetV1Adapter shapes: ((None, 32, 32, 3), (None,)), types: (tf.float32,  
tf.int64)>
```

6 Create callbacks

```
[69]: checkpoint_every_epoch = ModelCheckpoint(cp_path_every_epoch +  
        ↪ "checkpoint_{epoch}_{loss}_{val_loss}.ckpt",  
        save_best_only=False,  
        ↪ save_weights_only=True,  
        mode='auto', save_freq='epoch')  
checkpoint_best_epoch = ModelCheckpoint(cp_path_best_epoch + "checkpoint.ckpt",  
        save_best_only=True,  
        ↪ save_weights_only=True,  
        mode='auto', save_freq='epoch',  
        ↪ monitor="val_loss")  
early_stopping = EarlyStopping(monitor='val_loss', min_delta=0, patience=5,  
        restore_best_weights=True)
```

```

reduce_on_plateau = ReduceLROnPlateau(monitor='val_loss', factor=0.1,
    ↪patience=3)
term_NaN = TerminateOnNaN()
# lr_schedul = LearningRateScheduler(lambda epoch: 1e-4 * 10**(1 / (epoch + 1)))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    ↪histogram_freq=1)

```

```

[70]: callbacks = [checkpoint_every_epoch, checkpoint_best_epoch, early_stopping,
    reduce_on_plateau, term_NaN, tensorboard_callback]

```

7 Create MLP model

```

[71]: def create_MLP_model(img_shape):
    initializer = tf.keras.initializers.HeUniform()
    regularizer = tf.keras.regularizers.l2(1e-3)
    model = Sequential([
        Flatten(input_shape=img_shape),
        Dense(1024, activation='relu',
            kernel_regularizer=regularizer,
            ↪kernel_initializer=initializer),
        Dense(512, activation='relu',
            kernel_regularizer=regularizer,
            ↪kernel_initializer=initializer),
        Dense(256, activation='relu',
            kernel_regularizer=regularizer,
            ↪kernel_initializer=initializer),
        Dense(10, activation='softmax', name='Output_layer')
    ])
    model.summary()
    return model

model_mlp = create_MLP_model(img_shape)

```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
flatten_12 (Flatten)	(None, 3072)	0
dense_22 (Dense)	(None, 1024)	3146752
dense_23 (Dense)	(None, 512)	524800
dense_24 (Dense)	(None, 256)	131328
Output_layer (Dense)	(None, 10)	2570

```
=====
Total params: 3,805,450
Trainable params: 3,805,450
Non-trainable params: 0
-----
```

8 Compile and fit MLP model

```
[72]: model_mlp.compile(optimizer=tf.keras.optimizers.Adam(lr=0.0002),  
    ↪ loss='sparse_categorical_crossentropy', metrics=['acc'])  
history_mlp = model_mlp.fit(ds_train, epochs=50,  
    validation_data=ds_test, verbose=1,  
    callbacks=callbacks)  
history_mlp
```

Epoch 1/50

```
2/Unknown - 0s 125ms/step - loss: 6.1990 - acc:  
0.1289WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared  
to the batch time (batch time: 0.0241s vs `on_train_batch_end` time: 0.2217s).  
Check your callbacks.
```

```
WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the  
batch time (batch time: 0.0241s vs `on_train_batch_end` time: 0.2217s). Check  
your callbacks.
```

```
573/573 [=====] - 24s 42ms/step - loss: 3.7932 - acc:  
0.3320 - val_loss: 2.5923 - val_acc: 0.5119
```

Epoch 2/50

```
573/573 [=====] - 4s 6ms/step - loss: 2.1135 - acc:  
0.6092 - val_loss: 1.9413 - val_acc: 0.6204
```

Epoch 3/50

```
573/573 [=====] - 3s 6ms/step - loss: 1.6873 - acc:  
0.6763 - val_loss: 1.6817 - val_acc: 0.6607
```

Epoch 4/50

```
573/573 [=====] - 4s 6ms/step - loss: 1.4943 - acc:  
0.7082 - val_loss: 1.5095 - val_acc: 0.6978
```

Epoch 5/50

```
573/573 [=====] - 3s 6ms/step - loss: 1.3753 - acc:  
0.7293 - val_loss: 1.4175 - val_acc: 0.7108
```

Epoch 6/50

```
573/573 [=====] - 4s 6ms/step - loss: 1.2951 - acc:  
0.7419 - val_loss: 1.3519 - val_acc: 0.7232
```

Epoch 7/50

```
573/573 [=====] - 3s 6ms/step - loss: 1.2298 - acc:  
0.7536 - val_loss: 1.3061 - val_acc: 0.7294
```

Epoch 8/50

```
573/573 [=====] - 3s 6ms/step - loss: 1.1751 - acc:  
0.7631 - val_loss: 1.2537 - val_acc: 0.7394
```

Epoch 9/50
573/573 [=====] - 3s 6ms/step - loss: 1.1306 - acc:
0.7702 - val_loss: 1.2263 - val_acc: 0.7415
Epoch 10/50
573/573 [=====] - 3s 6ms/step - loss: 1.0908 - acc:
0.7765 - val_loss: 1.1936 - val_acc: 0.7465
Epoch 11/50
573/573 [=====] - 3s 6ms/step - loss: 1.0562 - acc:
0.7831 - val_loss: 1.1490 - val_acc: 0.7563
Epoch 12/50
573/573 [=====] - 3s 6ms/step - loss: 1.0252 - acc:
0.7873 - val_loss: 1.1260 - val_acc: 0.7599
Epoch 13/50
573/573 [=====] - 3s 6ms/step - loss: 0.9951 - acc:
0.7927 - val_loss: 1.0939 - val_acc: 0.7673
Epoch 14/50
573/573 [=====] - 4s 6ms/step - loss: 0.9689 - acc:
0.7990 - val_loss: 1.0727 - val_acc: 0.7724
Epoch 15/50
573/573 [=====] - 4s 6ms/step - loss: 0.9454 - acc:
0.8028 - val_loss: 1.0578 - val_acc: 0.7752
Epoch 16/50
573/573 [=====] - 4s 6ms/step - loss: 0.9211 - acc:
0.8088 - val_loss: 1.0234 - val_acc: 0.7846
Epoch 17/50
573/573 [=====] - 3s 6ms/step - loss: 0.9023 - acc:
0.8115 - val_loss: 1.0126 - val_acc: 0.7842
Epoch 18/50
573/573 [=====] - 3s 6ms/step - loss: 0.8874 - acc:
0.8143 - val_loss: 1.0164 - val_acc: 0.7812
Epoch 19/50
573/573 [=====] - 3s 6ms/step - loss: 0.8728 - acc:
0.8169 - val_loss: 1.0082 - val_acc: 0.7821
Epoch 20/50
573/573 [=====] - 3s 6ms/step - loss: 0.8582 - acc:
0.8193 - val_loss: 1.0003 - val_acc: 0.7830
Epoch 21/50
573/573 [=====] - 3s 6ms/step - loss: 0.8426 - acc:
0.8230 - val_loss: 0.9763 - val_acc: 0.7907
Epoch 22/50
573/573 [=====] - 3s 6ms/step - loss: 0.8308 - acc:
0.8259 - val_loss: 0.9807 - val_acc: 0.7875
Epoch 23/50
573/573 [=====] - 3s 6ms/step - loss: 0.8199 - acc:
0.8277 - val_loss: 0.9435 - val_acc: 0.7994
Epoch 24/50
573/573 [=====] - 3s 6ms/step - loss: 0.8117 - acc:
0.8291 - val_loss: 0.9456 - val_acc: 0.7976

Epoch 25/50
573/573 [=====] - 3s 6ms/step - loss: 0.8021 - acc:
0.8310 - val_loss: 0.9224 - val_acc: 0.8053
Epoch 26/50
573/573 [=====] - 3s 6ms/step - loss: 0.7884 - acc:
0.8347 - val_loss: 0.9288 - val_acc: 0.8009
Epoch 27/50
573/573 [=====] - 3s 6ms/step - loss: 0.7819 - acc:
0.8366 - val_loss: 0.9348 - val_acc: 0.7981
Epoch 28/50
573/573 [=====] - 3s 6ms/step - loss: 0.7733 - acc:
0.8387 - val_loss: 0.9373 - val_acc: 0.7964
Epoch 29/50
573/573 [=====] - 3s 6ms/step - loss: 0.6961 - acc:
0.8659 - val_loss: 0.8449 - val_acc: 0.8300
Epoch 30/50
573/573 [=====] - 3s 6ms/step - loss: 0.6868 - acc:
0.8691 - val_loss: 0.8419 - val_acc: 0.8307
Epoch 31/50
573/573 [=====] - 3s 6ms/step - loss: 0.6833 - acc:
0.8701 - val_loss: 0.8399 - val_acc: 0.8315
Epoch 32/50
573/573 [=====] - 3s 6ms/step - loss: 0.6802 - acc:
0.8707 - val_loss: 0.8369 - val_acc: 0.8316
Epoch 33/50
573/573 [=====] - 3s 6ms/step - loss: 0.6775 - acc:
0.8716 - val_loss: 0.8355 - val_acc: 0.8318
Epoch 34/50
573/573 [=====] - 4s 6ms/step - loss: 0.6748 - acc:
0.8721 - val_loss: 0.8324 - val_acc: 0.8323
Epoch 35/50
573/573 [=====] - 3s 6ms/step - loss: 0.6721 - acc:
0.8728 - val_loss: 0.8307 - val_acc: 0.8326
Epoch 36/50
573/573 [=====] - 3s 6ms/step - loss: 0.6697 - acc:
0.8735 - val_loss: 0.8285 - val_acc: 0.8332
Epoch 37/50
573/573 [=====] - 3s 6ms/step - loss: 0.6673 - acc:
0.8739 - val_loss: 0.8268 - val_acc: 0.8334
Epoch 38/50
573/573 [=====] - 3s 6ms/step - loss: 0.6651 - acc:
0.8742 - val_loss: 0.8250 - val_acc: 0.8336
Epoch 39/50
573/573 [=====] - 3s 6ms/step - loss: 0.6628 - acc:
0.8748 - val_loss: 0.8235 - val_acc: 0.8337
Epoch 40/50
573/573 [=====] - 3s 6ms/step - loss: 0.6607 - acc:
0.8752 - val_loss: 0.8219 - val_acc: 0.8343


```

Epoch 41/50
573/573 [=====] - 3s 6ms/step - loss: 0.6586 - acc:
0.8755 - val_loss: 0.8206 - val_acc: 0.8347
Epoch 42/50
573/573 [=====] - 3s 6ms/step - loss: 0.6566 - acc:
0.8761 - val_loss: 0.8195 - val_acc: 0.8346
Epoch 43/50
573/573 [=====] - 3s 6ms/step - loss: 0.6547 - acc:
0.8766 - val_loss: 0.8181 - val_acc: 0.8347
Epoch 44/50
573/573 [=====] - 3s 6ms/step - loss: 0.6528 - acc:
0.8771 - val_loss: 0.8167 - val_acc: 0.8352
Epoch 45/50
573/573 [=====] - 3s 6ms/step - loss: 0.6509 - acc:
0.8775 - val_loss: 0.8153 - val_acc: 0.8355
Epoch 46/50
573/573 [=====] - 3s 6ms/step - loss: 0.6491 - acc:
0.8781 - val_loss: 0.8141 - val_acc: 0.8355
Epoch 47/50
573/573 [=====] - 3s 6ms/step - loss: 0.6473 - acc:
0.8784 - val_loss: 0.8131 - val_acc: 0.8358
Epoch 48/50
573/573 [=====] - 3s 6ms/step - loss: 0.6456 - acc:
0.8787 - val_loss: 0.8121 - val_acc: 0.8362
Epoch 49/50
573/573 [=====] - 3s 6ms/step - loss: 0.6439 - acc:
0.8791 - val_loss: 0.8113 - val_acc: 0.8365
Epoch 50/50
573/573 [=====] - 3s 6ms/step - loss: 0.6422 - acc:
0.8797 - val_loss: 0.8101 - val_acc: 0.8367

```

[72]: <tensorflow.python.keras.callbacks.History at 0x7faebfd748d0>

9 Create CNN model

```

[73]: def create_CNN_model(img_shape):
        initializer = tf.keras.initializers.HeUniform()
        # regularizer dont improve; Dropout better
        regularizer = None # tf.keras.regularizers.l2(2e-5)

        model = Sequential([
                                Conv2D(32, 3, input_shape=(32, 32, 3), padding='same',
→activation='relu',
                                kernel_regularizer=regularizer,
→kernel_initializer=initializer,
                                name='Input_layer'),

```

```

        Conv2D(32, 3, padding='same', activation='relu',
              kernel_regularizer=regularizer,
↪kernel_initializer=initializer),
        BatchNormalization(),
        MaxPool2D(2),
        Dropout(0.4),
        Conv2D(64, 3, padding='same', activation='relu',
              kernel_regularizer=regularizer,
↪kernel_initializer=initializer),
        Conv2D(64, 3, padding='same', activation='relu',
              kernel_regularizer=regularizer,
↪kernel_initializer=initializer),
        BatchNormalization(),
        MaxPool2D(2),
        Dropout(0.4),
        Conv2D(128, 3, padding='same', activation='relu',
              kernel_regularizer=regularizer,
↪kernel_initializer=initializer),
        Conv2D(128, 3, padding='same', activation='relu',
              kernel_regularizer=regularizer,
↪kernel_initializer=initializer),
        BatchNormalization(),
        MaxPool2D(2),
        Dropout(0.4),
        Flatten(),
        #GlobalMaxPool2D(),
        Dense(128, activation='relu',
             kernel_regularizer=regularizer,
↪kernel_initializer=initializer),
        BatchNormalization(),
        Dropout(0.5),
        Dense(10, activation='softmax', name='Output_layer')
    ])
    model.summary()
    return model

```

```
model = create_CNN_model(img_shape)
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
Input_layer (Conv2D)	(None, 32, 32, 32)	896
conv2d_35 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_28 (Batc	(None, 32, 32, 32)	128

```

-----
max_pooling2d_21 (MaxPooling (None, 16, 16, 32)      0
-----
dropout_28 (Dropout)      (None, 16, 16, 32)      0
-----
conv2d_36 (Conv2D)      (None, 16, 16, 64)      18496
-----
conv2d_37 (Conv2D)      (None, 16, 16, 64)      36928
-----
batch_normalization_29 (Batc (None, 16, 16, 64)      256
-----
max_pooling2d_22 (MaxPooling (None, 8, 8, 64)      0
-----
dropout_29 (Dropout)      (None, 8, 8, 64)      0
-----
conv2d_38 (Conv2D)      (None, 8, 8, 128)      73856
-----
conv2d_39 (Conv2D)      (None, 8, 8, 128)      147584
-----
batch_normalization_30 (Batc (None, 8, 8, 128)      512
-----
max_pooling2d_23 (MaxPooling (None, 4, 4, 128)      0
-----
dropout_30 (Dropout)      (None, 4, 4, 128)      0
-----
flatten_13 (Flatten)      (None, 2048)      0
-----
dense_25 (Dense)      (None, 128)      262272
-----
batch_normalization_31 (Batc (None, 128)      512
-----
dropout_31 (Dropout)      (None, 128)      0
-----
Output_layer (Dense)      (None, 10)      1290
=====
Total params: 551,978
Trainable params: 551,274
Non-trainable params: 704
-----

```

10 Compile and fit CNN model

```

[74]: # clear previous history
!rm -rf checkpoint_best_epoch checkpoint_every_epoch logs_fit

model.compile(optimizer=tf.keras.optimizers.Adam(lr=0.001),
↳loss='sparse_categorical_crossentropy', metrics=['acc'])

```

```
history = model.fit(ds_train, epochs=50,
                    validation_data=ds_test, verbose=1,
                    callbacks=callbacks)
history
```

Epoch 1/50

2/Unknown - 0s 60ms/step - loss: 3.5918 - acc: 0.1133
WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0139s vs `on_train_batch_end` time: 0.1068s). Check your callbacks.

WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0139s vs `on_train_batch_end` time: 0.1068s). Check your callbacks.

573/573 [=====] - 9s 16ms/step - loss: 1.8782 - acc: 0.3980 - val_loss: 0.7196 - val_acc: 0.7725

Epoch 2/50

573/573 [=====] - 9s 15ms/step - loss: 0.6050 - acc: 0.8108 - val_loss: 0.3828 - val_acc: 0.8843

Epoch 3/50

573/573 [=====] - 9s 15ms/step - loss: 0.4561 - acc: 0.8602 - val_loss: 0.3222 - val_acc: 0.9049

Epoch 4/50

573/573 [=====] - 9s 15ms/step - loss: 0.3887 - acc: 0.8817 - val_loss: 0.2823 - val_acc: 0.9174

Epoch 5/50

573/573 [=====] - 9s 15ms/step - loss: 0.3530 - acc: 0.8935 - val_loss: 0.2882 - val_acc: 0.9146

Epoch 6/50

573/573 [=====] - 9s 15ms/step - loss: 0.3213 - acc: 0.9040 - val_loss: 0.2499 - val_acc: 0.9285

Epoch 7/50

573/573 [=====] - 9s 15ms/step - loss: 0.3012 - acc: 0.9096 - val_loss: 0.2308 - val_acc: 0.9345

Epoch 8/50

573/573 [=====] - 9s 15ms/step - loss: 0.2865 - acc: 0.9143 - val_loss: 0.2264 - val_acc: 0.9347

Epoch 9/50

573/573 [=====] - 9s 15ms/step - loss: 0.2691 - acc: 0.9196 - val_loss: 0.2074 - val_acc: 0.9414

Epoch 10/50

573/573 [=====] - 9s 15ms/step - loss: 0.2547 - acc: 0.9245 - val_loss: 0.2005 - val_acc: 0.9437

Epoch 11/50

573/573 [=====] - 9s 15ms/step - loss: 0.2455 - acc: 0.9282 - val_loss: 0.1920 - val_acc: 0.9467

Epoch 12/50

573/573 [=====] - 9s 15ms/step - loss: 0.2368 - acc:
0.9295 - val_loss: 0.1882 - val_acc: 0.9479
Epoch 13/50
573/573 [=====] - 9s 15ms/step - loss: 0.2291 - acc:
0.9320 - val_loss: 0.1943 - val_acc: 0.9446
Epoch 14/50
573/573 [=====] - 9s 15ms/step - loss: 0.2212 - acc:
0.9341 - val_loss: 0.1942 - val_acc: 0.9465
Epoch 15/50
573/573 [=====] - 9s 15ms/step - loss: 0.2115 - acc:
0.9381 - val_loss: 0.1797 - val_acc: 0.9503
Epoch 16/50
573/573 [=====] - 9s 15ms/step - loss: 0.2092 - acc:
0.9380 - val_loss: 0.1808 - val_acc: 0.9504
Epoch 17/50
573/573 [=====] - 9s 15ms/step - loss: 0.2029 - acc:
0.9400 - val_loss: 0.1795 - val_acc: 0.9521
Epoch 18/50
573/573 [=====] - 9s 15ms/step - loss: 0.1932 - acc:
0.9426 - val_loss: 0.1759 - val_acc: 0.9534
Epoch 19/50
573/573 [=====] - 9s 15ms/step - loss: 0.1872 - acc:
0.9449 - val_loss: 0.1701 - val_acc: 0.9551
Epoch 20/50
573/573 [=====] - 9s 15ms/step - loss: 0.1819 - acc:
0.9465 - val_loss: 0.1686 - val_acc: 0.9558
Epoch 21/50
573/573 [=====] - 9s 15ms/step - loss: 0.1814 - acc:
0.9470 - val_loss: 0.1689 - val_acc: 0.9557
Epoch 22/50
573/573 [=====] - 9s 15ms/step - loss: 0.1754 - acc:
0.9488 - val_loss: 0.1689 - val_acc: 0.9561
Epoch 23/50
573/573 [=====] - 9s 15ms/step - loss: 0.1699 - acc:
0.9496 - val_loss: 0.1662 - val_acc: 0.9566
Epoch 24/50
573/573 [=====] - 9s 15ms/step - loss: 0.1652 - acc:
0.9509 - val_loss: 0.1698 - val_acc: 0.9552
Epoch 25/50
573/573 [=====] - 9s 15ms/step - loss: 0.1624 - acc:
0.9518 - val_loss: 0.1672 - val_acc: 0.9560
Epoch 26/50
573/573 [=====] - 9s 15ms/step - loss: 0.1610 - acc:
0.9532 - val_loss: 0.1764 - val_acc: 0.9538
Epoch 27/50
573/573 [=====] - 9s 15ms/step - loss: 0.1389 - acc:
0.9589 - val_loss: 0.1626 - val_acc: 0.9597
Epoch 28/50

```

573/573 [=====] - 9s 15ms/step - loss: 0.1337 - acc:
0.9603 - val_loss: 0.1617 - val_acc: 0.9595
Epoch 29/50
573/573 [=====] - 9s 15ms/step - loss: 0.1279 - acc:
0.9621 - val_loss: 0.1614 - val_acc: 0.9589
Epoch 30/50
573/573 [=====] - 9s 15ms/step - loss: 0.1248 - acc:
0.9626 - val_loss: 0.1621 - val_acc: 0.9586
Epoch 31/50
573/573 [=====] - 9s 15ms/step - loss: 0.1184 - acc:
0.9651 - val_loss: 0.1608 - val_acc: 0.9599
Epoch 32/50
573/573 [=====] - 9s 15ms/step - loss: 0.1180 - acc:
0.9656 - val_loss: 0.1625 - val_acc: 0.9589
Epoch 33/50
573/573 [=====] - 9s 15ms/step - loss: 0.1202 - acc:
0.9641 - val_loss: 0.1602 - val_acc: 0.9597
Epoch 34/50
573/573 [=====] - 9s 15ms/step - loss: 0.1161 - acc:
0.9655 - val_loss: 0.1605 - val_acc: 0.9589
Epoch 35/50
573/573 [=====] - 9s 15ms/step - loss: 0.1150 - acc:
0.9661 - val_loss: 0.1594 - val_acc: 0.9601
Epoch 36/50
573/573 [=====] - 9s 15ms/step - loss: 0.1149 - acc:
0.9656 - val_loss: 0.1605 - val_acc: 0.9606
Epoch 37/50
573/573 [=====] - 9s 15ms/step - loss: 0.1116 - acc:
0.9665 - val_loss: 0.1610 - val_acc: 0.9599
Epoch 38/50
573/573 [=====] - 9s 15ms/step - loss: 0.1093 - acc:
0.9676 - val_loss: 0.1633 - val_acc: 0.9590
Epoch 39/50
573/573 [=====] - 9s 15ms/step - loss: 0.1067 - acc:
0.9689 - val_loss: 0.1600 - val_acc: 0.9609
Epoch 40/50
573/573 [=====] - 9s 15ms/step - loss: 0.1077 - acc:
0.9678 - val_loss: 0.1598 - val_acc: 0.9609

```

[74]: <tensorflow.python.keras.callbacks.History at 0x7fae5ddd57b8>

11 Evaluate model

```
[75]: model.evaluate(ds_test)
```

```

204/204 [=====] - 1s 5ms/step - loss: 0.1594 - acc:
0.9601

```

```
[75]: [0.15944360196590424, 0.9601259827613831]
```

```
[81]: predictions = model.predict(ds_test)[:10]


for images, labels in ds_test.take(1):
    labels = labels.numpy()
    images = images.numpy()
    x = np.array([i for i in range(0, 10)])

    print("MLP classifier")
    predictions = model_mlp.predict(ds_test)[:10]
    plt.figure(figsize=(20, 6))
    for (ii, val) in enumerate(predictions):
        plt.subplot(2, 5, ii + 1)
        plt.bar(x, val)
        plt.title(labels[ii])
    plt.show()



    print("CNN classifier")
    predictions = model.predict(ds_test)[:10]
    plt.figure(figsize=(20, 6))
    for (ii, val) in enumerate(predictions):
        plt.subplot(2, 5, ii + 1)
        plt.bar(x, val)
        plt.title(labels[ii])
    plt.show()

    plt.figure(figsize=(10, 3.5))
    for (ii, image) in enumerate(images[:10]):
        plt.subplot(2, 5, ii + 1)
        plt.imshow(image)
        plt.xticks([])
        plt.yticks([])
    plt.show()
```

MLP classifier



CNN classifier



12 Visualize history

```
[77]: %tensorboard --logdir logs_fit
```

<IPython.core.display.Javascript object>

```
[78]: def show_hository(history):
    plt.figure(figsize=(12, 3))
    # print(history.history.keys())
    plt.subplot(1, 3, 1)
    plt.plot(history.history['acc'])
    plt.plot(history.history['val_acc'])
    plt.legend(['acc', 'val_acc'])
    plt.subplot(1, 3, 2)
```




```

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend(['loss', 'val_loss'])
plt.subplot(1, 3, 3)
plt.plot(history.history['lr'])
plt.legend(['lr'])
plt.show()


print('MLP history')
show_history(history_mlp)
print('CNN history')
show_history(history)

```

MLP history



CNN history



13 Load weights and save model

```

[79]: model.load_weights(tf.train.latest_checkpoint(cp_path_best_epoch))
      model.evaluate(ds_test)

      model.save('final_model')

```

```
# keras model
model.save('final_tf_model.h5')
```

```
204/204 [=====] - 1s 5ms/step - loss: 0.1594 - acc:
0.9601
```

```
INFO:tensorflow:Assets written to: final_model/assets
```

```
INFO:tensorflow:Assets written to: final_model/assets
```

```
[80]: !ls
```

```
checkpoint_best_epoch  final_model          logs_fit
checkpoint_every_epoch  final_tf_model.h5    sample_data
```